

VizRec 2.0: Towards a holistic workflow for visual data exploration

Doris Jung-Lin Lee, Aditya Parameswaran
{jlee782,adityagp}@illinois.edu
University of Illinois, Urbana-Champaign

Abstract

test abstract

1 Introduction

Visualization — insights.

1. Visual analysis help discover insights, etc.
2. supporting cycle of visual data exploration.
3. In this paper, we introduce various challenges in the — visual analysis, and discuss — solutions —.
4. Challenge #1: Precise search
5. finding the right vis and relevant data is hard,
6. we discuss *zenvisage* as one use case and solution in this space (Section 2)
7. Challenge #2: Need for in-the-loop support. how do people query visually?
8. During our PD + others, hypothesis + in the loop considerations is important. We discuss relevant work and our findings in Section 3.
9. Need to formulate complex expressive queries, language is good, but need interface (e.g visual primitives) to support this. (Section 4)
10. top-down, bottom-up — point to need for bottom-up recommendations (Section 5)
11. Challenge #3: Vague and intelligent search
12. Users might not always have something they want to start with, supporting vague querying (Section 4)
13. Challenge #4: Cold-start recommendation for data understanding. One aspect of this is to gain understanding of dataset, (e.g. representative and outliers in *zenvisage*, bottom up exploration)we discuss STORYBOARD as one example solution addressing this problem (Section 5).

Copyright 0000 IEEE. Personal use of this material is permitted. However, permission to reprint/republish this material for advertising or promotional purposes or for creating new collective works for resale or redistribution to servers or lists, or to reuse any copyrighted component of this work in other works must be obtained from the IEEE.

Bulletin of the IEEE Computer Society Technical Committee on Data Engineering

2 Precise Visual Querying

Visual analysis often reveal important anomalies or trends in the data[?]. However, it is often challenging to find the appropriate piece of information to realize these insights.

2.1 Motivating Example

Astronomers from the The Dark Energy Survey (DES)[?] are interested in finding anomalous time series to discover astrophysical transients (objects whose brightness changes dramatically as a function of time), such as supernova explosions or quasars. When trying to find celestial objects corresponding to supernovae, which have a specific pattern of brightness over time, scientists need to individually inspect the visualizations of each object until they find ones that match the pattern. With more than 400 million objects in their catalog, each having their own set of time series brightness measurement, the process of manually exploring a large number of visualizations is not only error-prone, but also overwhelming for scientists who do not have extensive knowledge about their dataset.

The astronomy use case highlights a common challenge in exploratory data analysis (EDA). There is often a large space of possible visualizations that could be generated from a given dataset and manual search through this large collection is inefficient. Visualization authoring tools such as Tableau and Excel focusses on presenting one visualization at a time. There is no systematic way to create, compare, and query large collections of visualizations.

2.2 Effortless Data Exploration with *zenvisage*

The challenges presented earlier points to a need for a tool that enables users to create and search through large collections of visualizations. We developed *zenvisage* a visual query system that . *zenvisage* is built on top of a querying language called ZQL, which provides a mechanism for managing collections of visualizations[?]. Contrary prior work on visualization languages for specifying visual encodings of individual visualizations [?, ?], ZQL supports high-level queries over visualizations, such as composing, sorting, filtering a collection of visualization. Example functionalities of ZQL includes:

- Comparing across a collection of visualizations by iterating over one or more x, y, z attributes while fixing other attributes (e.g. Find me all the cities where the housing prices starts off high then becomes lower over time. Here we vary along CITY while keeping X=TIME,Y=AVG(PRICE) fixed.)
- Finding the top-k visualizations whose y values are most or least similar from a queried visualization.
- Finding a pair of X and Y axes where the visualizations for two specific products ‘stapler’ and ‘chair’ differ the most

Given a ZQL query, *zenvisage* parses it into a graph of visual component nodes(contains visualization information, such as X, Y columns) and task nodes (common and user-defined primitives for processing visual components, such as sort-filter). *zenvisage* then performs query optimization to merge together multiple nodes, as well as reducing the processing time required for individual visualization components. Using the optimized query plan, the executor compiles visual nodes into SQL queries for retrieving the visualization data and post-processes the result via the defined operations.

While ZQL provides powerful mechanism for expressively specifying queries on large collections visualizations, writing ZQL queries can be daunting for novice users. Therefore, we extracted a typical workflow of visualization querying (finding top-k most similar visualization from a collection with fixed X,Y while varying Z) to allow users to formulate ZQL queries through interactions (such as drag-and-drop, sketching), while allowing them to construct ZQL queries via a frontend table input. *zenvisage* maps the frontend interactions

into ZQL queries submitted to the backend and renders the query results as a ranked list of visualizations in the Results panel as shown in Figure 4.2. *zenvisage* is a full-fledged visual querying system that supports a variety of querying interactions as described in Figure 4.2. In the following section, we will discuss the design process of how we developed this visual query system and the lessons that we have learned for designing future visual data exploration systems.

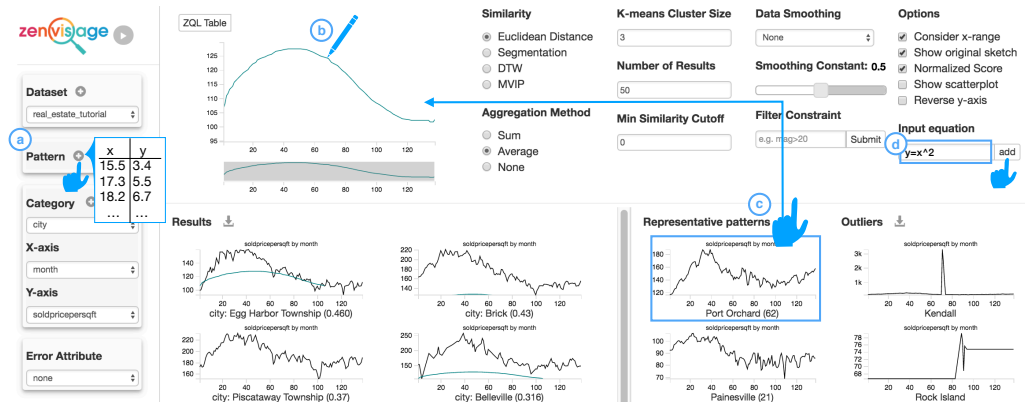


Figure 1: *zenvisage* offers a variety of querying modalities, including: a) uploading a sample pattern from an external dataset as a query, b) sketching a query pattern, c) dragging-and-dropping an existing pattern from the dataset, and d) inputting an equation as a query.

3 Hypothesis Formation and the cycle of visual analysis

3.1 Visual Querying in the framework of Data Sensemaking

In developing *zenvisage*, we collaborated with scientists from astronomy, genetics, and material science in a year-long participatory design process [?]. In particular, we study how various features impacts rapid visual querying and analysis, and how a visual query system (VQS) like *zenvisage* fit into a scientists’ analysis workflow. Our findings offered design guidelines for improving the usability and adoption of next-generation VQSs. More importantly, in this paper, we focus on applying our findings on visual querying to the context of supporting the full cycle of visual data exploration. We highlight two of the key findings related to this goal below:

Top-down and Bottom-up Querying Modalities We employed Pirolli and Card’s [?] information foraging framework for domain-experts to contextualize our study results. Pirolli and Card’s notional model distinguishes between information processing tasks that are *top-down* (from theory to data) and *bottom-up* (from data to theory). In the context of visual querying, users employ top-down approaches by starting with a hypothesis on what patterns to look for and express it through sketching or inputting an equation (Figure 4.2b,d). On the other hand, bottom-up approaches originate from the data (or equivalently, the visualization). For example, the user may drag and drop a visualization of interest in the dataset as the input query or upload a visualization from an external dataset (Figure 4.2a,c).

Our interactions with the scientists showed that different modalities for inputting a query can be useful for different problem contexts. To our surprise, despite the prevalence of sketch-to-query systems in literature, only two out of our nine users had a practical usage for querying by sketching. Overall, we found that bottom-up querying via drag-and-drop was more intuitive and more commonly used than top-down querying methods, such as sketching or input equations.

One of the main reason why participants did not find sketching useful was that they often do not start their analysis with a pattern in mind. Later, their intuition about what to query is derived from other visualizations that they see in the VQS, in which case it made more sense to query using those visualizations as examples directly.

While —, formulating top-down queries in a ad-hoc manner can be challenging —. This — why participants did not find top-down querying through inputting equations and sketching —. *Ad-hoc, sketch-only model for visualization querying is insufficient without* . Despite functional fitting being common in scientific data analysis, Figure ?? shows that querying by equation is also unpopular for similar reasons. In addition, the visualizations for both astronomy and biology exhibit complex processes that could not be written down as an equation analytically. However, even when analytical relationships do exist (in the case of material science), it is challenging to formulate functional forms in an prescriptive, ad-hoc manner.

While the usage of each querying feature may vary from one participant to the next, a key design principle that came from this finding was the need for visual query systems to provide visualization recommendations that can help analysts jumpstart their exploration. Our results indicate that *bottom-up querying approaches are preferred over top-down when the users have no desired patterns in mind*, which is commonly the case for exploratory data analysis. For example, the analyst may examine the top-most similar visualizations that lie in a cluster and finding visualizations that are similar to an object of interest that exhibits a desired pattern.

Our *zenvisage* work — Bottom up and top down querying in VQS facilitates rapid insight discovery.

Construction of multi-step queries workflow integration , complex queries, While our visual interface —, —it was unable to capture all the —. [Give some example queries that didn't work]. Balance between language/querying complexity versus expressiveness. More importantly pointed towards a need for vague querying. Give some examples of vague querying.

3.2 Supporting cycle of visual data exploration

Our participatory design findings points towards future — in supporting a cycle of —. —advocate —cycle of visual analysis.

- Essential ingredient in facilitating intelligent vague querying and exploration.
- This is a human process ([?, ?])
- Iterative Hypothesis Exploration/Refinement : argue that the following properties is important to sustain this “cycle of visual analysis”

3.3 Towards 3‘I’s of rapid hypothesis generation support

Given our observations from the participatory design study, we distill several desiderata for the next generation VQSS. Towards 3‘I’s Interactive, Iterative, Informative (Give examples from the ZV-TVCG paper)

Interactive flow: (how natural is it to move between analysis steps, facilitate fluid analysis and not get “stuck”) : interactivity, feedback (latter is quite unexplored), and recommendation, expressivity (how easy is it to express what to do via interactions) and diversity of actions that could be performed.

Iterative: query refinement, dialogue (not a one-shot query) Joining the flow: Section 4 focuses on the first two items .

Informative: not just task-based interestingness but more explanation-based (causality, introduce distribution awareness notion in viz-sum), focussed on data understanding, which we will discuss in Section 5

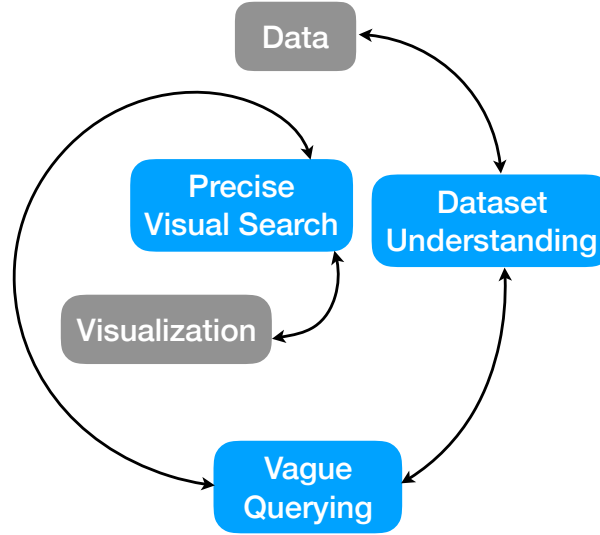


Figure 2: Cycle of visual data exploration.

3.4 Challenges Ahead

The goal here is to help novice submit precise queries without SQL background, easy to use interface. Our study found that VQS does more than just this, but still not enough.

- Precise Search Fail to understand intricacies of user need/intent, need more expressivity/flexibility for querying.
- No perfect training workload, real-world data + task is noisy and complex.
- towards more holistic model for insight discovery

4 Vague Intelligent Search

Accounting for user interaction, mental models. More global objective taking into account user with the goal of dataset understanding rather than task completion.

4.1 Challenges

- Inferring user intent in querying and context is important (both in terms of user input and what is recommended)
- tools can not assume user has querying intention. exploration without intention, user don't know what they are searching for - Recommendation.
- The important thing here is identifying what should be done by the system v.s. requested from user. Inappropriate choice of these will result in lack of expressibility and user feeling lack of control of analysis, limiting exploration.
- Need for a unified framework of inference to take all of these into account (e.g. natural language, etc)

4.2 STORYBOARD: Navigating Through Data Slices with Hierarchical Summary of Visualizations

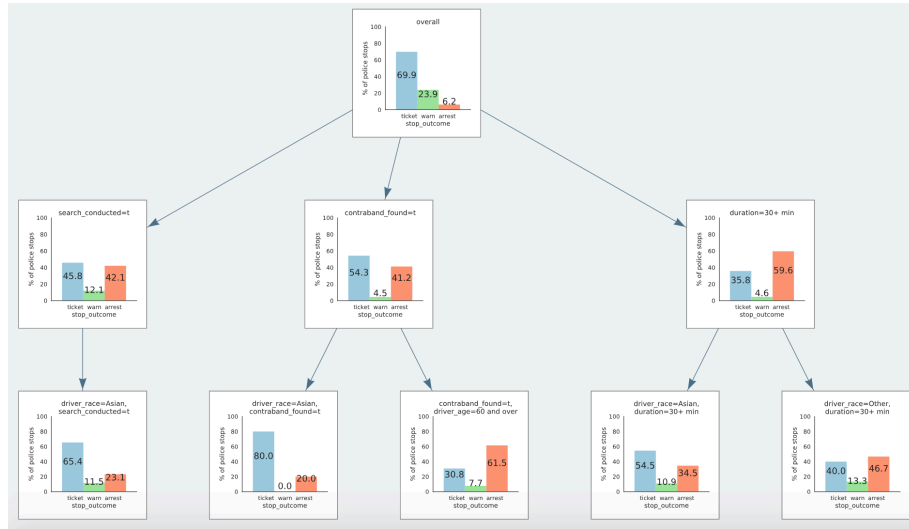


Figure 3: Example dashboard with generated from the Police Stop Dataset [?], which contains records of police stops that resulted in a warning, ticket, or an arrest. The attributes in the dataset include driver gender, age, race, and the stop time of day, whether a search was conducted, and whether contraband was found. We generate a dashboard of visualizations with bar charts with x-axis as the stop outcome (whether the police stop resulted in a ticket, warning, or arrest/summons) and y-axis as the percentage of police stops that led to this outcome.

Data is agnostic to the user, intention —, by building tools—, Section 2 to 4 have focussed on extracting what user want from data. bridging together what user want from data, what data has to offer, supporting interactive discourse between the two. Either using one-size-fits-all statistics, templates, heuristics as a solution or problem only applicable to a subset of analytic tasks[?, ?].

5 Towards Dataset Understanding

5.1 Challenges

- Problem of cold-start recommendation (as discussed earlier use may not always know what to query for)
- related works have focussed on making specification easier , but not really trying to understand user intent or what might the user want to see.
- Within a dataset, structure and provenance is essential to help users navigate and provide users with sense of coverage and completion. This is an important but underexplored area. (viz-sum, Sarvghad et al 2017)
- schema and attribute understanding (coverage, etc)

5.2 Examples

- understanding distributions (distribution awareness)
- providing overview recommendations (representative trends and outliers)

6 Concluding Remarks